

Relational Algebra

CS 4400

Relational Algebra

Relational model specifies structures and constraints, relational algebra provides retrieval operations

- Formal foundation for relational model operations
- Basis for internal query optimization in RDBMS
- Parts of relational algebra found in SQL

Basic Rules

- Relational algebra expressions operate on relations and produce relations as results
- Relational algebra expressions can be chained

SELECT

$\sigma_{\langle condition \rangle}(R)$

- R is the name of a relation
- $\langle condition \rangle$ is a boolean condition on the values of attributes in the tuples of R

The select operation returns all the tuples from R for which $\langle condition \rangle$ is true.

SELECT Example

Given the following data for **pet**:

```
+-----+---+-----+-----+
| shelter_id | id | name  | breed  |
+-----+---+-----+-----+
```

| | | | |
|---|---|--------|-----------|
| 1 | 1 | Chloe | Mix |
| 1 | 2 | Dante | GSD |
| 1 | 3 | Heidi | Dachshund |
| 2 | 1 | Bailey | Mix |
| 2 | 2 | Sophie | Lab |
| 2 | 3 | Heidi | Dachshund |

$\sigma_{\text{breed}='mix'}(\text{pet})$ returns:

| shelter_id | id | name | breed |
|------------|----|--------|-------|
| 1 | 1 | Chloe | Mix |
| 2 | 1 | Bailey | Mix |

Properties of SELECT

- Result of $\sigma_{\langle \text{condition} \rangle}(R)$ has same schema as R , i.e., same attributes
- SELECT is commutative, e.g.,

$$\sigma_{\langle c1 \rangle}(\sigma_{\langle c2 \rangle}(R)) = \sigma_{\langle c2 \rangle}(\sigma_{\langle c1 \rangle}(R))$$
- Cascaded SELECTs can be replaced by single SELECT with conjunction of conditions, e.g.

$$\sigma_{\langle c1 \rangle}(\sigma_{\langle c2 \rangle}(R)) = \sigma_{\langle c1 \rangle \text{ AND } \langle c2 \rangle}(R)$$
- Result of $\sigma_{\langle \text{condition} \rangle}(R)$ has equal or fewer tuples than R

PROJECT

$\pi_{\langle \text{attributelist} \rangle}(R)$

- R is the name of a relation
- $\langle \text{attributelist} \rangle$ is a subset of the attributes of relation R

The project operation returns all the tuples in R but with only the attributes in $\langle \text{attribute} - \text{list} \rangle$

PROJECT Example

| shelter_id | id | name | breed |
|------------|----|------|-------|
|------------|----|------|-------|

| | | | | | |
|--|---|---|--------|-----------|--|
| | 1 | 1 | Chloe | Mix | |
| | 1 | 2 | Dante | GSD | |
| | 1 | 3 | Heidi | Dachshund | |
| | 2 | 1 | Bailey | Mix | |
| | 2 | 2 | Sophie | Lab | |
| | 2 | 3 | Heidi | Dachshund | |

$\pi_{name,breed}(pet) =$

| | | | |
|--|--------|-----------|--|
| | name | breed | |
| | Chloe | Mix | |
| | Dante | GSD | |
| | Heidi | Dachshund | |
| | Bailey | Mix | |
| | Sophie | Lab | |

Notice that the duplicate tuple $\langle Heidi, Dachshund \rangle$ was removed. Results of relational algebra operations are sets.

Properties of PROJECT

- Number of tuples returned by PROJECT is less than or equal to the number of tuples in the input relation because result is a set, i.e., $|\pi_{\langle attrs \rangle}(R)| \leq |R|$
 - What if $\langle attrs \rangle$ includes a key of R ?
- PROJECT is *not* commutative. In fact $\pi_{\langle attrs1 \rangle}(\pi_{\langle attrs2 \rangle}(R))$ is only a correct expression if $\langle attrs2 \rangle$ contains the attributes in $\langle attrs1 \rangle$. In this case the result is simply $\pi_{\langle attrs1 \rangle}(R)$.

Combining PROJECT and SELECT

| | | | | | |
|--|------------|----|--------|-----------|--|
| | shelter_id | id | name | breed | |
| | 1 | 1 | Chloe | Mix | |
| | 1 | 2 | Dante | GSD | |
| | 1 | 3 | Heidi | Dachshund | |
| | 2 | 1 | Bailey | Mix | |
| | 2 | 2 | Sophie | Lab | |
| | 2 | 3 | Heidi | Dachshund | |

```

+-----+-----+-----+
πname(σbreed='Mix'(pet)) produces:
+-----+
| name  |
+-----+
| Chloe |
| Bailey|
+-----+

```

Intermediate Results

Previous *in-line expression* could be split up into multiple steps with named intermediate results.

```

πname(σbreed='Mix'(pet))
becomes:
MIXES ← σbreed='Mix'(pet)
RESULT ← πname(MIXES)

```

RENAME

- Rename relation R to S :
 $\rho_S(R)$
- Rename attributes of R to B_1, \dots, B_n :
 $\rho_{(B_1, \dots, B_n)}(R)$
- Rename R to S and attributes to B_1, \dots, B_n :
 $\rho_{S(B_1, \dots, B_n)}(R)$

Binary Operators

- UNION, $R \cup S$, is set of all tuples in either R or S
- INTERSECTION, $R \cap S$, is set of all tuples in both R and S
- SET DIFFERENCE, $R - S$, is set of all tuples in R but not in S

Operands must be *union compatible*, or *type compatible*. For R and S to be union compatible:

- Degree of R must be same as degree of S
- For each attribute A_i in R and B_i in S , $dom(A_i) = dom(B_i)$

Cartesian Product

$R \times S$ Creates “super-tuples” by concatenating every tuple in R with every tuple in S .

$$R(A_1, \dots, A_n) \times S(B_1, \dots, B_m) = Q(A_1, \dots, A_n, B_1, \dots, B_m)$$

Notice that

- Q has degree $n + m$
- $|q(Q)| = |r(R)| \times |s(S)|$

Note that the book abuses notation a bit and writes that last bullet as $|Q| = |R| \times |S|$

Cartesian Product Example

shelter

| id | name |
|----|---------|
| 1 | Howell |
| 2 | Mansell |

pet

| shelter_id | id | name | breed |
|------------|----|--------|-----------|
| 1 | 1 | Chloe | Mix |
| 1 | 2 | Dante | GSD |
| 1 | 3 | Heidi | Dachshund |
| 2 | 1 | Bailey | Mix |
| 2 | 2 | Sophie | Lab |
| 2 | 3 | Heidi | Dachshund |

Cross Product Example

| sid | sname | shelter_id | pid | pname | breed |
|-----|---------|------------|-----|-------|-------|
| 1 | Howell | 1 | 1 | Chloe | Mix |
| 2 | Mansell | 1 | 1 | Chloe | Mix |
| 1 | Howell | 1 | 2 | Dante | GSD |

| | | | | | | | | | | | | |
|---|---|--|---------|--|---|--|---|--|--------|--|-----------|--|
| | 2 | | Mansell | | 1 | | 2 | | Dante | | GSD | |
| | 1 | | Howell | | 1 | | 3 | | Heidi | | Dachshund | |
| | 2 | | Mansell | | 1 | | 3 | | Heidi | | Dachshund | |
| | 1 | | Howell | | 2 | | 1 | | Bailey | | Mix | |
| | 2 | | Mansell | | 2 | | 1 | | Bailey | | Mix | |
| | 1 | | Howell | | 2 | | 2 | | Sophie | | Lab | |
| | 2 | | Mansell | | 2 | | 2 | | Sophie | | Lab | |
| | 1 | | Howell | | 2 | | 3 | | Heidi | | Dachshund | |
| | 2 | | Mansell | | 2 | | 3 | | Heidi | | Dachshund | |
| +-----+-----+-----+-----+-----+-----+-----+ | | | | | | | | | | | | |

Note that we've also done a **RENAME** to disambiguate names and ids:

$\rho_{(sid,sname,shelter_id,pid,pname,breed)}(shelter \times pet)$

Cross Product and Select

Cross product meaningful when combined with **SELECT**.

$\sigma_{sid=shelter_id}(\rho_{(sid,sname,shelter_id,pid,pname,breed)}(shelter \times pet))$

| | | | | | | | | | | | | |
|---|-----|--|---------|--|------------|--|-----|--|--------|--|-----------|--|
| | sid | | sname | | shelter_id | | pid | | pname | | breed | |
| | 1 | | Howell | | 1 | | 1 | | Chloe | | Mix | |
| | 1 | | Howell | | 1 | | 2 | | Dante | | GSD | |
| | 1 | | Howell | | 1 | | 3 | | Heidi | | Dachshund | |
| | 2 | | Mansell | | 2 | | 1 | | Bailey | | Mix | |
| | 2 | | Mansell | | 2 | | 2 | | Sophie | | Lab | |
| | 2 | | Mansell | | 2 | | 3 | | Heidi | | Dachshund | |
| +-----+-----+-----+-----+-----+-----+-----+ | | | | | | | | | | | | |

$CROSSED \leftarrow shelter \times pet$

$RENAMED \leftarrow \rho_{(sid,sname,shelter_id,pid,pname,breed)}(CROSSED)$

$RESULT \leftarrow \sigma_{sid=shelter_id}(RENAMED)$

Join

JOIN is a **CARTESIAN PRODUCT** followed by **SELECT**

$R \bowtie_{<joincondition>} S$

Where

- R and S are relations
- $< joincondition >$ is a boolean condition on values of tuples from R and S

$R \bowtie_{\langle joincondition \rangle} S$ returns the tuples in $R \times S$ that satisfy the $\langle joincondition \rangle$

Join Conditions

$\langle joincondition \rangle$ is of the form $A_i \theta B_j$

- A_i is an attribute of R , B_j is an attribute of S
- $dom(A_i) = dom(B_j)$

A $\langle joincondition \rangle$ can be a conjunction of simple conditions, e.g.:

$\langle c_1 \rangle AND \langle c_2 \rangle ... AND \langle c_n \rangle$

Join Example

worker

| id | name | supervisor_id | shelter_id |
|----|--------|---------------|------------|
| 1 | Tom | NULL | 1 |
| 2 | Jie | 1 | 1 |
| 3 | Ravi | 2 | 1 |
| 4 | Alice | 2 | 1 |
| 5 | Aparna | NULL | 2 |
| 6 | Bob | 5 | 2 |
| 7 | Xaoxi | 6 | 2 |
| 8 | Rohan | 6 | 2 |

shelter

| id | name |
|----|---------|
| 1 | Howell |
| 2 | Mansell |

Join Example

$worker \bowtie_{shelter_id=sid} \rho_{(sid,sname)}(shelter)$

| id | name | supervisor_id | shelter_id | sid | sname |
|----|--------|---------------|------------|-----|---------|
| 1 | Tom | NULL | 1 | 1 | Howell |
| 2 | Jie | 1 | 1 | 1 | Howell |
| 3 | Ravi | 2 | 1 | 1 | Howell |
| 4 | Alice | 2 | 1 | 1 | Howell |
| 5 | Aparna | NULL | 2 | 2 | Mansell |
| 6 | Bob | 5 | 2 | 2 | Mansell |
| 7 | Xaoxi | 6 | 2 | 2 | Mansell |
| 8 | Rohan | 6 | 2 | 2 | Mansell |

Notice that we had to use renaming of attributes in *shelter*.

A join operation in which the comparison operator θ is $=$ is called an *equijoin*.

Natural Join

Notice that the `shelter_id` attribute was repeated in the previous equijoin result. A **NATURAL JOIN** is a equijoin in which the redundant attribute has been removed.

$R * S$

Where

- R and S have an attribute with the same name and same domain which is automatically chosen as the equijoin attribute

Natural Join Example

Recall the first join example. If we rename the `id` attribute to `shelter_id` we can use a natural join:

$\rho_{(shelter_id, sname)}(shelter) * worker$

| shelter_id | sname | id | name | supervisor_id |
|------------|---------|----|--------|---------------|
| 1 | Howell | 1 | Tom | NULL |
| 1 | Howell | 2 | Jie | 1 |
| 1 | Howell | 3 | Ravi | 2 |
| 1 | Howell | 4 | Alice | 2 |
| 2 | Mansell | 5 | Aparna | NULL |
| 2 | Mansell | 6 | Bob | 5 |
| 2 | Mansell | 7 | Xaoxi | 6 |

| | | | | | | | | |
|---------|---|---------|--|---------|-------|---------|---|---------|
| | 2 | Mansell | | 8 | Rohan | | 6 | |
| +-----+ | | +-----+ | | +-----+ | | +-----+ | | +-----+ |

Outer Joins

The joins we've discussed so far have been inner joins. Result relations of inner joins include only tuples from the joined tables that match the join condition.

Outer join results include tuples that matched, and tuples that didn't match the join condition.

Left Outer Join

$R_{<joincondition>}S$

Where

- R and S are relations
- $< joincondition >$ is a boolean condition on values of tuples from R and S

$R_{<joincondition>}S$ returns the tuples in $R \times S$ that satisfy the $< joincondition >$ as well as the tuples from R that don't match the join condition. In the result relation the unmatched tuples from R are null-padded to give them the correct degree in the result.

Left Outer Join Example

author

| | | | |
|-----------|------------|-----------|--|
| +-----+ | +-----+ | +-----+ | |
| author_id | first_name | last_name | |
| +-----+ | +-----+ | +-----+ | |
| 1 | John | McCarthy | |
| 2 | Dennis | Ritchie | |
| 3 | Ken | Thompson | |
| 4 | Claude | Shannon | |
| 5 | Alan | Turing | |
| 6 | Alonzo | Church | |
| 7 | Perry | White | |
| 8 | Moshe | Vardi | |
| 9 | Roy | Batty | |
| +-----+ | +-----+ | +-----+ | |

book

| book_id | book_title | month | year | editor |
|---------|------------|----------|------|--------|
| 1 | CACM | April | 1960 | 8 |
| 2 | CACM | July | 1974 | 8 |
| 3 | BST | July | 1948 | 2 |
| 4 | LMS | November | 1936 | 7 |
| 5 | Mind | October | 1950 | NULL |
| 6 | AMS | Month | 1941 | NULL |
| 7 | AAAI | July | 2012 | 9 |
| 8 | NIPS | July | 2012 | 9 |

Return all the authors. For all the authors who are editors, show their books.

Authors and Edited Books

Show all the authors. For all the authors who are editors, show their books.

$R_{author_i d=editor} S$

| author_id | first_name | last_name | book_id | book_title | month | year | editor |
|-----------|------------|-----------|---------|------------|----------|------|--------|
| 8 | Moshe | Vardi | 1 | CACM | April | 1960 | 8 |
| 8 | Moshe | Vardi | 2 | CACM | July | 1974 | 8 |
| 2 | Dennis | Ritchie | 3 | BST | July | 1948 | 2 |
| 7 | Perry | White | 4 | LMS | November | 1936 | 7 |
| 9 | Roy | Batty | 7 | AAAI | July | 2012 | 9 |
| 9 | Roy | Batty | 8 | NIPS | July | 2012 | 9 |
| 1 | John | McCarthy | NULL | NULL | NULL | NULL | NULL |
| 3 | Ken | Thompson | NULL | NULL | NULL | NULL | NULL |
| 4 | Claude | Shannon | NULL | NULL | NULL | NULL | NULL |
| 5 | Alan | Turing | NULL | NULL | NULL | NULL | NULL |
| 6 | Alonzo | Church | NULL | NULL | NULL | NULL | NULL |

Notice how attribute values are padded to the right in a left outer join.

Right Outer Join

$R_{<joincondition>} S$

Where

- R and S are relations

- $\langle joincondition \rangle$ is a boolean condition on values of tuples from R and S

$R_{\langle joincondition \rangle} S$ returns the tuples in $R \times S$ that satisfy the $\langle joincondition \rangle$ as well as the tuples from S that don't match the join condition. In the result relation the unmatched tuples from S are null-padded to give them the correct degree in the result.

Right Outer Join Example

Show all the books. For books with editors, show their editors.

$R_{author_id=editor} S$

| author_id | first_name | last_name | book_id | book_title | month | year | editor |
|-----------|------------|-----------|---------|------------|----------|------|--------|
| 8 | Moshe | Vardi | 1 | CACM | April | 1960 | 8 |
| 8 | Moshe | Vardi | 2 | CACM | July | 1974 | 8 |
| 2 | Dennis | Ritchie | 3 | BST | July | 1948 | 2 |
| 7 | Perry | White | 4 | LMS | November | 1936 | 7 |
| NULL | NULL | NULL | 5 | Mind | October | 1950 | NULL |
| NULL | NULL | NULL | 6 | AMS | Month | 1941 | NULL |
| 9 | Roy | Batty | 7 | AAAI | July | 2012 | 9 |
| 9 | Roy | Batty | 8 | NIPS | July | 2012 | 9 |

Notice how attribute values are padded to the left in a right outer join.

Full Outer Join

$R_{\langle joincondition \rangle} S$

Where

- R and S are relations
- $\langle joincondition \rangle$ is a boolean condition on values of tuples from R and S

$R_{\langle joincondition \rangle} S$ returns the tuples in $R \times S$ that satisfy the $\langle joincondition \rangle$ as well as the tuples from both R and S that don't match the join condition. In the result relation the unmatched tuples are null-padded to give them the correct degree in the result.

Full Outer Join Example

Show all authors and books, matching editors with their books.

$R_{author_d=editor}S$

| author_id | first_name | last_name | book_id | book_title | month | year | editor |
|-----------|------------|-----------|---------|------------|----------|------|--------|
| 8 | Moshe | Vardi | 1 | CACM | April | 1960 | 8 |
| 8 | Moshe | Vardi | 2 | CACM | July | 1974 | 8 |
| 2 | Dennis | Ritchie | 3 | BST | July | 1948 | 2 |
| 7 | Perry | White | 4 | LMS | November | 1936 | 7 |
| 9 | Roy | Batty | 7 | AAAI | July | 2012 | 9 |
| 9 | Roy | Batty | 8 | NIPS | July | 2012 | 9 |
| 1 | John | McCarthy | NULL | NULL | NULL | NULL | NULL |
| 3 | Ken | Thompson | NULL | NULL | NULL | NULL | NULL |
| 4 | Claude | Shannon | NULL | NULL | NULL | NULL | NULL |
| 5 | Alan | Turing | NULL | NULL | NULL | NULL | NULL |
| 6 | Alonzo | Church | NULL | NULL | NULL | NULL | NULL |
| NULL | NULL | NULL | 5 | Mind | October | 1950 | NULL |
| NULL | NULL | NULL | 6 | AMS | Month | 1941 | NULL |

Review Question 1

Given the $r(book)$:

| book_id | book_title | month | year | editor |
|---------|------------|----------|------|--------|
| 1 | CACM | April | 1960 | 8 |
| 2 | CACM | July | 1974 | 8 |
| 3 | BST | July | 1948 | 2 |
| 4 | LMS | November | 1936 | 7 |
| 5 | Mind | October | 1950 | 7 |
| 6 | AMS | Month | 1941 | 7 |
| 7 | AAAI | July | 2012 | 9 |
| 8 | NIPS | July | 2012 | 9 |

How many tuples are in $\pi_{book_title}(book)$? (Notice I've abused notation here, since that's how you'll see it in the book and on the test.)

Review Question 1 Answer

7:

| book_title |
|------------|
| CACM |
| BST |
| LMS |
| Mind |
| AMS |
| AAAI |
| NIPS |

The *book_title* appears twice in *book* and the result of a relational algebra expression is a set.

Review Question 2

Given the relation $r(book)$:

| book_id | book_title | month | year | editor |
|---------|------------|----------|------|--------|
| 1 | CACM | April | 1960 | 8 |
| 2 | CACM | July | 1974 | 8 |
| 3 | BST | July | 1948 | 2 |
| 4 | LMS | November | 1936 | 7 |
| 5 | Mind | October | 1950 | 7 |
| 6 | AMS | Month | 1941 | 7 |
| 7 | AAAI | July | 2012 | 9 |
| 8 | NIPS | July | 2012 | 9 |

Which books were published before 1960 or after 2000?

Review Question 2

Which books were published before 1960 or after 2000?

$\sigma_{year < 1960}(book) \cup \sigma_{year > 2000}(book)$

| book_id | book_title | month | year | editor |
|---------|------------|-------|------|--------|
|---------|------------|-------|------|--------|

| | | | | |
|---|------|----------|------|---|
| 3 | BST | July | 1948 | 2 |
| 4 | LMS | November | 1936 | 7 |
| 5 | Mind | October | 1950 | 7 |
| 6 | AMS | Month | 1941 | 7 |
| 7 | AAAI | July | 2012 | 9 |
| 8 | NIPS | July | 2012 | 9 |

Review Question 3

Given:

worker

| id | name | supervisor_id | shelter_id |
|----|--------|---------------|------------|
| 1 | Tom | NULL | 1 |
| 2 | Jie | 1 | 1 |
| 3 | Ravi | 2 | 1 |
| 4 | Alice | 2 | 1 |
| 5 | Aparna | NULL | 2 |
| 6 | Bob | 5 | 2 |
| 7 | Xaoxi | 6 | 2 |
| 8 | Rohan | 6 | 2 |

shelter

| id | name |
|----|---------|
| 1 | Howell |
| 2 | Mansell |

How would we find the names of all the workers who work at Mansell?

Review Question 3 Answer

How would we find all the workers who work at Mansell?

$SHELTERS \leftarrow \rho_{sid, sname}(shelter)$

$MANSELLERS \leftarrow$

$worker \bowtie_{shelter_id=sid \wedge sname='Mansell'} (SHELTERS)$

Gives:

| id | name | supervisor_id | shelter_id | sid | sname |
|----|--------|---------------|------------|-----|---------|
| 5 | Aparna | NULL | 2 | 2 | Mansell |
| 6 | Bob | 5 | 2 | 2 | Mansell |
| 7 | Xaoxi | 6 | 2 | 2 | Mansell |
| 8 | Rohan | 6 | 2 | 2 | Mansell |

then ...

Review Question 3 Answer

$\pi_{name}(MANSELLERS)$

gives:

| name |
|--------|
| Aparna |
| Bob |
| Xaoxi |
| Rohan |

Full inline expression:

$\pi_{name}(worker \bowtie_{shelter_id=sid \wedge sname='Mansell'} \rho_{(sid,sname)}(shelter))$