

Relational Data Model

Christopher Simpkins

Relational Data Model

A **relation schema** $R(A_1, \dots, A_n)$ is a relation name R and a list of attributes A_1, \dots, A_n .

Each attribute A_i is the name of a role played by some domain D .

- ▶ Example: $AUTHOR(author_id, first_name, last_name)$
 - ▶ $dom(A_1)$ (or $dom(author_id)$) is integer

A **database schema** is a collection of relation schemas.

- ▶ Example: $PUBS$ database has relation schemas $BOOK$, $AUTHOR$, and PUB (for publication, not public house)

Relations and Databases

A **relation**, or **relation state**, $r(R)$ is a **set** of tuples that conform to a **relation schema** R .

- ▶ Example: $r(AUTHOR) =$

| author__id | first__name | last__name |
|------------|-------------|------------|
| 1 | John | McCarthy |
| 4 | Claude | Shannon |
| 5 | Alan | Turing |
| 6 | Alonzo | Church |

A **database** is a set of relations.

Tuples

A **tuple** is an **ordered list** of values that is part of a relation

- ▶ Example: $t_1 = \langle 1, 'John', 'McCarthy' \rangle$

Each value in the tuple is that tuple's value for the corresponding attribute of the relation schema.

Example: (these are equivalent notations):

- ▶ $t_1[first_name] = 'John'$ (bracket notation)
- ▶ $t_1.first_name = 'John'$ (object notation)
- ▶ $t_1[2] = 'John'$ (positional notation)

The **degree** or **arity** of a relation schema is the number of attributes it has.

- ▶ Example: *AUTHOR* has degree 3.

An Example Relation

Relation Name

STUDENT

Attributes

Tuples

| Name | Ssn | Home_phone | Address | Office_phone | Age | Gpa |
|----------------|-------------|------------|----------------------|--------------|-----|------|
| Benjamin Bayer | 305-61-2435 | 373-1616 | 2918 Bluebonnet Lane | NULL | 19 | 3.21 |
| Chung-cha Kim | 381-62-1245 | 375-4409 | 125 Kirby Road | NULL | 18 | 2.89 |
| Dick Davidson | 422-11-2320 | NULL | 3452 Elgin Road | 749-1253 | 25 | 3.53 |
| Rohan Panchal | 489-22-1100 | 376-9821 | 265 Lark Lane | 749-6492 | 28 | 3.93 |
| Barbara Benson | 533-69-1238 | 839-8461 | 7384 Fontana Lane | NULL | 19 | 3.25 |

Figure 5.1

The attributes and tuples of a relation STUDENT.

Attributes and Domains

Each attribute has a name and a **domain**

- ▶ The name describes the role played by the attribute
 - ▶ Example: the *first_name* attribute of the *AUTHOR* schema plays the role of the first name of an author represented by a tuple in a $r(AUTHOR)$ relation.
- ▶ The domain is a set of atomic values that a tuple may have for that attribute.
- ▶ A **logical definition** of a domain specifies a simple type such as integer or string, and a **data type** or **format**
 - ▶ Example: USA_phone_number as $(ddd)ddd - dddd$, where d is a digit

Mathematical Definition of Relation

Given $R(A_1, \dots, A_n)$,

$$\blacktriangleright r(R) \subseteq (dom(A_1) \times dom(A_2) \times \dots \times dom(A_n))$$

The total number of values, or **cardinality**, of a domain D is $|D|$.

So the maximum number of tuples that could possibly be in $r(R)$ is

$$\blacktriangleright |dom(A_1)| * |dom(A_2)| * \dots * |dom(A_n)|$$

Properties of Relations

- ▶ Atomicity of values, i.e., the First Normal Form assumption
 - ▶ Attribute values in tuples are indivisible, e.g., no compound or multivalued attributes as in EER models
- ▶ Nulls
 - ▶ Unknown, not applicable, not existing
- ▶ Closed world assumption
 - ▶ Facts not asserted explicitly are assumed to be false

Kinds of Constraints

- ▶ Inherent model-based (or **implicit**) constraints
 - ▶ domain constraints, atomic attribute values
- ▶ Schema-based (or **explicit**) constraints
 - ▶ keys, referential integrity
- ▶ Application-based (or semantic constraints), a.k.a., business rules

Superkeys

A **superkey** SK is a set of attributes of a relation schema R such that

$$t_i[SK] \neq t_j[SK]$$

for any $i \neq j$.

In other words, the values of the superkey attributes of a tuple uniquely identify the tuple within the relation.

By the definition of the relational model, the full attribute set of a relation schema is a **default superkey**.

Keys

A **minimal superkey** is a superkey removing an attribute would make it no longer unique, and thus no longer a superkey.

We call a minimal superkey a **key**.

A relation schema may have several keys. We call these **candidate keys** and choose one arbitrarily to be the **primary key**.

We underline the primary key in a relation schema.

- ▶ Example: *AUTHOR*(*author_id*, *first_name*, *last_name*)

Database Integrity Constraints

- ▶ Domain constraints - Attribute values in tuples must be in domain for that attribute
- ▶ Key constraints - No two tuples can have the same values for the primary key
- ▶ Entity Integrity Constraints - No tuple can have a NULL value for its primary key attribute
- ▶ Referential Integrity Constraints - Tuples in one relation referencing tuples in another relation
- ▶ Semantic Integrity Constraints - Constraints on values of attributes that cannot be specified in the databases DDL

Referential Integrity Constraints

A foreign key value from a tuple in one relation must refer to nothing, or to the primary key for an existing tuple in another relation. Formally:

Given relation schemas R_1 and R_2 , a set of attributes FK in R_1 is a foreign key referencing R_2 if

- ▶ the attributes in FK in R_1 have same domains as PK in R_2
- ▶ Given some t_1 in $r_1(R_1)$ and t_2 in $r_2(R_2)$, either $t_1[FK] = t_2[PK]$ or $t_1[FK]$ is NULL.

R_1 is the referencing relation, R_2 is the referenced relation.

Diagramming FK Relationships

EMPLOYEE

| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|

DEPARTMENT

| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|-------|----------------|---------|----------------|
|-------|----------------|---------|----------------|

DEPT_LOCATIONS

| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|
|----------------|------------------|

PROJECT

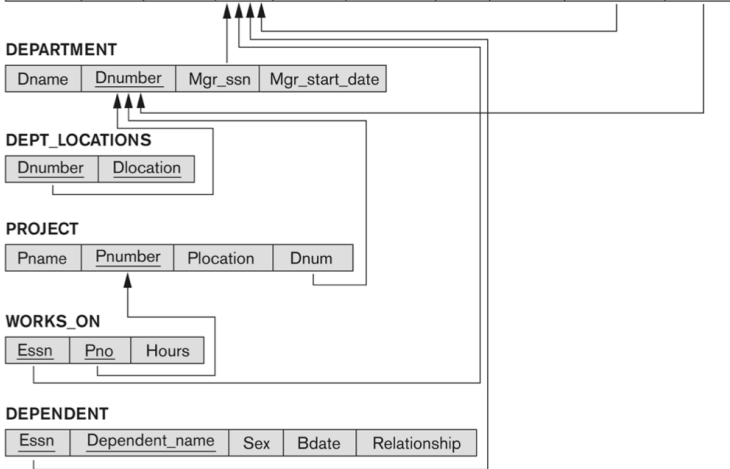
| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-------|----------------|-----------|------|
|-------|----------------|-----------|------|

WORKS_ON

| <u>Essn</u> | <u>Pno</u> | Hours |
|-------------|------------|-------|
|-------------|------------|-------|

DEPENDENT

| <u>Essn</u> | <u>Dependent_name</u> | Sex | Bdate | Relationship |
|-------------|-----------------------|-----|-------|--------------|
|-------------|-----------------------|-----|-------|--------------|



Semantic Integrity Constraints

- ▶ Can't be specified in DDL
- ▶ Can be checked with triggers and assertions
- ▶ Usually checked in application code

Example: salary of an employee cannot exceed the salary of the employee's supervisor.

Constraint Violations on Insert

- ▶ Domain constraints
 - ▶ Insert a tuple with an attribute value not in attribute's domain
- ▶ Key constraints
 - ▶ Insert a tuple with a key that's already in the relation state
- ▶ Entity integrity constraints
 - ▶ Insert a tuple with a NULL value for any part of the primary key
- ▶ Referential integrity constraints
 - ▶ Insert a tuple in a referring relation whose FK does not appear as a PK value in any tuple of the referenced relation

Constraint Violations on Update

- ▶ Domain constraints
 - ▶ Update a tuple with an attribute value not in attribute's domain
- ▶ Key constraints
 - ▶ Update a tuple with a key value that already appears in another tuple in the relation
- ▶ Entity integrity constraints
 - ▶ Update a tuple with a NULL value for any part of the primary key
- ▶ Referential integrity constraints
 - ▶ Update a tuple in a referring relation with a FK does not appear as a PK value in any tuple of the referenced relation

Constraint Violations on Delete

- ▶ Referential integrity

Employee - Department Example

EMPLOYEE

| Fname | Minit | Lname | Ssn | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

DEPARTMENT

| Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|----------------|---------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |