

# Relational Data Model

Christopher Simpkins

# Relational Data Model

A **relation schema**  $R(A_a, \dots, A_n)$  is a relation name  $R$  and a list of attributes  $A_1, \dots, A_n$ .

Each attribute  $A_i$  is the name of a role played by some domain  $D$ .

- ▶ Example:  $AUTHOR(author\_id, first\_name, last\_name)$ 
  - ▶  $dom(A_1)$  (or  $dom(author\_id)$ ) is integer

A **database schema** is a collection of relation schemas.

- ▶ Example:  $PUBS$  database has relation schemas  $BOOK$ ,  $AUTHOR$ , and  $PUB$  (for publication, not public house)

# Relations and Databases


A **relation**, or **relation state**,  $r(R)$  is a **set** of tuples that conform to a **relation schema**  $R$ .

- ▶ Example:  $r(AUTHOR) =$

author_id	first_name	last_name
1	John	McCarthy
4	Claude	Shannon
5	Alan	Turing
6	Alonzo	Church

A **database** is a set of relations.

# An Example Relation



**Figure 5.1**

The attributes and tuples of a relation STUDENT.

# Tuples

A **tuple** is an **ordered list** of values that is part of a relation

- ▶ Example:  $t_1 = \langle 1, 'John', 'McCarthy' \rangle$

Each value in the tuple is that tuple's value for the corresponding attribute of the relation schema.

Example: (these are equivalent notations):

- ▶  $t_1[first\_name] = "John"$  (bracket notation)
- ▶  $t_1.first\_name = "John"$  (object notation)
- ▶  $t_1[2] = "John"$  (positional notation)

The **degree** or **arity** of a relation schema is the number of attributes it has.

- ▶ Example: *AUTHOR* has degree 3.

# Attributes and Domains

Each attribute has a name and a **domain**

- ▶ The name describes the role played by the attribute
  - ▶ Example: the *first\_name* attribute of the *AUTHOR* schema plays the role of the first name of an author represented by a tuple in a  $r(AUTHOR)$  relation.
- ▶ The domain is a set of atomic values that a tuple may have for that attribute.
- ▶ A **logical definition** of a domain specifies a simple type such as integer or string, and a **data type** or **format**
  - ▶ Example: 'USA<sub>phonenumber</sub>' as  $(ddd)ddd - dddd$ , where  $d$  is a digit

# Mathematical Definition of Relation

Given  $R(A_1, \dots, A_n)$ ,

$$\blacktriangleright r(R) \subseteq (\text{dom}(A_1) \times \text{dom}(A_2) \times \dots \times \text{dom}(A_n))$$

The total number of values, or **cardinality**, of a domain  $D$  is  $|D|$ .

So the maximum number of tuples that could possibly be in  $r(R)$  is

$$\blacktriangleright |\text{dom}(A_1)| \times |\text{dom}(A_2)| \times \dots \times |\text{dom}(A_n)|$$

# Properties of Relations

- ▶ Atomicity of values, i.e., the First Normal Form assumption
  - ▶ Attribute values in tuples are indivisible, e.g., no compound or multivalued attributes as in EER models
- ▶ Nulls
  - ▶ Unknown, not applicable, not existing
- ▶ Closed world assumption
  - ▶ Facts not asserted explicitly are assumed to be false



# Constraints

- ▶ Inherent model-based (or **implicit**) constraints
  - ▶ domain constraints, atomic attribute values
- ▶ Schema-based (or **explicit**) constraints
  - ▶ keys, referential integrity
- ▶ Application-based (or semantic constraints), a.k.a., business rules

# Superkeys

A **superkey**  $SK$  is a set of attributes of a relation schema  $R$  such that

$$t_i[SK] \neq t_j[SK]$$

for any  $i \neq j$ .

In other words, the values of the superkey attributes of a tuple uniquely identify the tuple within the relation.

By the definition of the relational model, the full attribute set of a relation schema is a **default superkey**.

# Keys

A **minimal superkey** is a superkey removing an attribute would make it no longer unique, and thus no longer a superkey.

We call a minimal superkey a **key**.

A relation schema may have several keys. We call these **candidate keys** and choose one arbitrarily to be the **primary key**.

We underline the primary key in a relation schema.

- ▶ Example: *AUTHOR*(*author\_id*, *first\_name*, *last\_name*)

# Database Integrity Constraints

- ▶ Domain constraints - Attribute values in tuples must be in domain for that attribute
- ▶ Key constraints - No two tuples can have the same values for the primary key
- ▶ Entity Integrity Constraints - No tuple can have a NULL value for its primary key attribute
- ▶ Referential Integrity Constraints - Tuples in one relation referencing tuples in another relation
- ▶ Semantic Integrity Constraints - Constraints on values of attributes that cannot be specified in the databases DDL

# Referential Integrity Constraints

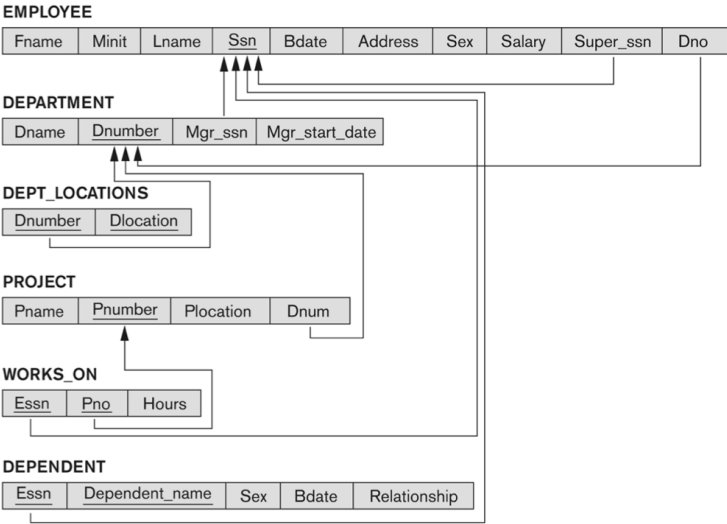
A foreign key value from a tuple in one relation must refer to nothing, or to the primary key for an existing tuple in another relation. Formally:

Given relation schemas  $R_1$  and  $R_2$ , a set of attributes  $FK$  in  $R_1$  is a foreign key referencing  $R_2$  if

- ▶ the attributes in  $FK$  in  $R_1$  have same domains as  $PK$  in  $R_2$
- ▶ Given some  $t_1$  in  $r_1(R_1)$  and  $t_2$  in  $r_2(R_2)$ , either  $t_1[FK] = t_2[PK]$  or  $t_1[FK]$  is NULL.

$R_1$  is the referencing relation,  $R_2$  is the referenced relation.

# Diagramming FK Relationships



# Semantic Integrity Constraints

- ▶ Can't be specified in DDL
- ▶ Can be checked with triggers and assertions
- ▶ Usually checked in application code

Example: salary of an employee cannot exceed the salary of the employee's supervisor.

# Constraint Violations on Insert

- ▶ Domain constraints
  - ▶ Insert a tuple with an attribute value not in attribute's domain
- ▶ Key constraints
  - ▶ Insert a tuple with a key that's already in the relation state
- ▶ Entity integrity constraints
  - ▶ Insert a tuple with a NULL value for any part of the primary key
- ▶ Referential integrity constraints
  - ▶ Insert a tuple in a referring relation whose FK does not appear as a PK value in any tuple of the referenced relation



# Constraint Violations on Update

- ▶ Domain constraints
  - ▶ Update a tuple with an attribute value not in attribute's domain
- ▶ Key constraints
  - ▶ Update a tuple with a key value that already appears in another tuple in the relation
- ▶ Entity integrity constraints
  - ▶ Update a tuple with a NULL value for any part of the primary key
- ▶ Referential integrity constraints
  - ▶ Update a tuple in a referring relation with a FK does not appear as a PK value in any tuple of the referenced relation

# Constraint Violations on Delete

- ▶ Referential integrity

# Employee - Department Example

**EMPLOYEE**

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

**DEPARTMENT**

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

**DEPT\_LOCATIONS**

Dnumber	Dlocation
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston