

# Advanced SQL

# NULL

The special value NULL could mean:

- ▶ Unknown
- ▶ Unavailable
- ▶ Not Applicable

# Three-Valued Logic - AND

AND	TRUE	FALSE	UNKNOWN	———	———	———	———
TRUE	TRUE	FALSE	UNKNOWN	FALSE	FALSE	FALSE	
FALSE	UNKNOWN	UNKNOWN	FALSE	UNKNOWN			



# Three-Valued Logic - NOT

NOT		TRUE		——— ———		TRUE		FALSE		FALSE		TRUE	
		UNKNOWN		UNKNOWN									

# Comparisons with NULL Values

Each NULL is distinct, so comparisons with  $<$ ,  $>$ , and  $=$  don't make sense.

To compare with null, use SQL operator IS, e.g., "Which books don't have editors?":

```
SELECT * FROM book WHERE editor IS NULL;
```

Inner joins include only tuples for which the join condition evaluates to TRUE.

# Queries in Depth

```
SELECT [DISTINCT] <select_header> FROM <source_tables>  
WHERE <filter_expression>  
GROUP BY <grouping_expressions>  
HAVING <filter_expression>  
ORDER BY <ordering_expressions>  
LIMIT <count> OFFSET <count>
```

- ▶ The table is the fundamental data abstraction in a relational database.
- ▶ The select command returns its result as a table
- ▶ Think of a select statement as creating a pipeline, each stage of which produces an intermediate working table

# The SELECT Pipeline

The evaluation order of select clauses is approximately:

1. FROM <source\_tables> - Designates one or more source tables and combines them together into one large working table.

1. WHERE <filter\_expression> - Filters specific rows of working table
2. GROUP BY <grouping\_expressions> - Groups sets of rows in the working table based on similar values
3. SELECT <select\_heading> - Defines the result set columns and (if applicable) grouping aggregates.
4. HAVING <filter\_expression> - Filters specific rows of the grouped table. Requires a GROUP BY
5. DISTINCT - Eliminates duplicate rows.
6. ~ORDER BY <ordering\_expressions> - Sorts the rows of the result set
7. OFFSET <count> - Skips over rows at the beginning of the result set. Requires a LIMIT.
8. LIMIT <count> - Limits the result set output to a specific number of rows.



# SELECT Pipeline at a Glance

1. FROM <source\_tables>
2. WHERE <filter\_expression>
3. GROUP BY <grouping\_expressions>
4. SELECT <select\_heading>
5. HAVING <filter\_expression>
6. DISTINCT
7. ORDER BY <ordering\_expressions>
8. OFFSET <count>
9. LIMIT <count>

Evaluation order determines what can be cross referenced in clauses.

# Aggregate Functions

Operate on groups of rows. Some common ones: COUNT, SUM, AVG

```
mysql> select count(*) from book;
```

```
+-----+  
| count(*) |  
+-----+  
|          8 |  
+-----+
```

There are 8 rows in the book table.

```
mysql> select count(editor) from book;
```

```
+-----+  
| count(editor) |  
+-----+  
|                6 |  
+-----+
```

Notice that COUNT doesn't count NULL values.



# Aggregate Functions on Groups

Aggregate functions apply some function to the rows grouped together by a GROUP BY clause.

How many papers did each author write?

```
mysql> select author_id, last_name, count(author_id)
-> from author join author_pub using (author_id)
-> join pub using (pub_id)
-> group by author_id;
```

author_id	last_name	count(author_id)
1	McCarthy	1
2	Ritchie	1
3	Thompson	1
4	Shannon	1
5	Turing	2
6	Church	1

Aggregate function is applied to column in GROUP BY.

# Sorting, Aliasing, and Limiting

Who wrote the most publications?

```
mysql> select author_id, last_name, count(author_id) as  
        pub_count  
        -> from author join author_pub using (author_id) join  
        pub using (pub_id)  
        -> group by author_id  
        -> order by pub_count desc;
```

```
+-----+-----+-----+  
| author_id | last_name | pub_count |  
+-----+-----+-----+  
|          5 | Turing   |          2 |  
|          1 | McCarthy |          1 |  
|          2 | Ritchie  |          1 |  
|          6 | Church   |          1 |  
|          3 | Thompson |          1 |  
|          4 | Shannon  |          1 |  
+-----+-----+-----+  
6 rows in set (0.00 sec)
```

Notice that we also used an alias so we could reference the count in the ORDER BY clause

# Limiting Results

If we want only the answer from the last query we can use LIMIT:  
Who wrote the most publications?

```
mysql> select author_id, last_name, count(author_id) as  
        pub_count  
        -> from author join author_pub using (author_id) join  
            pub using (pub_id)  
        -> group by author_id  
        -> order by pub_count desc  
        -> limit 1;
```

```
+-----+-----+-----+  
| author_id | last_name | pub_count |  
+-----+-----+-----+  
|          5 | Turing   |          2 |  
+-----+-----+-----+  
1 row in set (0.00 sec)
```

## HAVING

In the previous query we got the top author by pub count. If we want all authors having a particular pub count, we can use a HAVING clause.

```
mysql> select author_id, last_name, count(author_id) as  
        pub_count  
        -> from author join author_pub using (author_id)  
        ->   join pub using (pub_id)  
        -> group by author_id  
        -> having pub_count = 1;
```

Author_id	last_name	pub_count
1	McCarthy	1
2	Ritchie	1
3	Thompson	1
4	Shannon	1
6	Church	1

We can use comparisons like  $<$ ,  $>$ . Notice that Turing is not in the result.

# HAVING vs. WHERE Conditions

Functionally HAVING and WHERE do the same thing: they filter-in tuples. The difference is where they are evaluated in the SELECT pipeline.

- ▶ WHERE is evaluated only after the FROM clause that selects the source tables, so WHERE clauses can only reference expressions that do not contain aggregate functions
- ▶ HAVING is evaluated after GROUP BY, and SELECT, so HAVING clauses can reference any result column

Be aware that rows filtered out by a WHERE clause will not be included in a GROUP BY clause.



## WHERE vs. HAVING Example

WHERE clause can't refer to column aliases and aggregates in the SELECT list or apply functions to groups created by GROUP BY clauses.

```
mysql> select author_id, last_name, count(author_id) as
      pub_count
      -> from author natural join author_pub natural join
      pub
      -> where pub_count = 1
      -> group by author_id;
ERROR 1054 (42S22): Unknown column 'pub_count' in 'where
      clause'
```

HAVING can refer to select columns.

```
mysql> select author_id, last_name, count(author_id) as
      pub_count
      -> from author natural join author_pub natural join
      pub
      -> group by author_id
      -> having pub_count = 1;
```

```
+-----+-----+-----+
| author_id | last_name | pub_count |
+-----+-----+-----+
```

# The IN Operator

```
mysql> select * from book where month in ('April',  
      'July');
```

book_id	book_title	month	year	editor
1	CACM	April	1960	8
2	CACM	July	1974	8
3	BST	July	1948	2
7	AAAI	July	2012	9
8	NIPS	July	2012	9

5 rows in set (0.00 sec)

## Nested Queries, a.k.a., Sub-Selects

List all the books published in the same month in which an issue of CACM was published.

```
mysql> select book_title, month
-> from book
-> where month in (select month
->                  from book
                  where book_title = 'CACM');
```

```
+-----+-----+
| book_title | month |
+-----+-----+
| CACM      | April |
| CACM      | July  |
| BST       | July  |
| AAI       | July  |
| NIPS      | July  |
+-----+-----+
5 rows in set (0.00 sec)
```

# Simple Summation

Here are the data in the dorm table:

```
mysql> select * from dorm;
+-----+-----+-----+
| dorm_id | name      | spaces |
+-----+-----+-----+
|      1 | Armstrong |    124 |
|      2 | Brown     |    158 |
|      3 | Caldwell  |    158 |
+-----+-----+-----+
3 rows in set (0.00 sec)
```

What is the total capacity (number of spaces) for all dorms?

# SUM

To find the total capacity for all dorms, sum the spaces column:

```
mysql> select sum(spaces) from dorm;
```

```
+-----+
```

```
| sum(spaces) |
```

```
+-----+
```

```
|          440 |
```

```
+-----+
```

```
1 row in set (0.00 sec)
```

Or use a column alias in the select list to make output clearer:

```
mysql> select sum(spaces) as total_capacity from dorm;
```

```
+-----+
```

```
| total_capacity |
```

```
+-----+
```

```
|          440 |
```

```
+-----+
```

```
1 row in set (0.00 sec)
```

## Grouping and Counting

What is the occupancy of each dorm?

First, get a feel for the data:

```
mysql> select * from dorm join student using (dorm_id)
      order by dorm.name;
```

dorm_id	name	spaces	student_id	name	gpa
1	Armstrong	124	1	Alice	3.60
1	Armstrong	124	2	Bob	2.70
1	Armstrong	124	3	Cheng	3.90
2	Brown	158	4	Dhruv	3.40
2	Brown	158	5	Ellie	4.00
2	Brown	158	6	Fong	2.30
3	Caldwell	158	7	Gerd	4.00
3	Caldwell	158	8	Hal	2.20
3	Caldwell	158	9	Isaac	2.00
3	Caldwell	158	10	Jacque	4.00

We can see that there are three groups of dorms in the result, which we could group by `dorm_id` or `dorm.name`.

# Dorm Occupancy

So we group by dorm.name and count the rows in each group.

```
mysql> select dorm.name as dorm_name, count(*) as  
        occupancy  
        -> from dorm join student using (dorm_id)  
        -> group by dorm.name;
```

```
+-----+-----+  
| dorm_name | occupancy |  
+-----+-----+  
| Armstrong |         3 |  
| Brown     |         3 |  
| Caldwell  |         4 |  
+-----+-----+  
3 rows in set (0.00 sec)
```

# Ordering

```
mysql> select dorm.name as dorm_name, count(*) as  
        occupancy  
        -> from dorm join student using (dorm_id)  
        -> group by dorm.name  
        -> order by occupancy desc;
```

```
+-----+-----+  
| dorm_name | occupancy |  
+-----+-----+  
| Caldwell |         4 |  
| Armstrong |        3 |  
| Brown    |        3 |  
+-----+-----+  
3 rows in set (0.00 sec)
```



# Nested Queries

Which dorms have fewer occupants than Caldwell?

Step 1: how many occupants in Caldwell?

```
mysql> select count(*) as caldwell_occupancy  
-> from dorm join student using(dorm_id)  
-> where dorm.name = 'caldwell';
```

```
+-----+  
| caldwell_occupancy |  
+-----+  
|                  4 |  
+-----+  
1 row in set (0.00 sec)
```

## Occupancy Less than Caldwell

Now we use the previous "caldwell<sub>occupancy</sub>" query as a subquery.

```
mysql> select dorm.name as dorm_name, count(*) as
        occupancy
        -> from dorm join student using (dorm_id)
        -> group by dorm_name
        -> having occupancy < (select count(*) as
            caldwell_occupancy
        ->                               from dorm join student
            using(dorm_id)
        ->                               where dorm.name = 'caldwell');
```

```
+-----+-----+
| dorm_name | occupancy |
+-----+-----+
| Armstrong |         3 |
| Brown     |         3 |
+-----+-----+
2 rows in set (0.00 sec)
```

Notice that we couldn't use a where clause here because occupancy is computed from a group, which isn't available at the WHERE stage of the SQL SELECT pipeline.

## Extended Example: Which dorm has the highest average GPA?

- ▶ Step 1: Group students and their GPAs by dorm.
- ▶ Step 2: Get the average GPAs of each dorm.
- ▶ Step 3: Get the max avg GPA from step 2.

## Step 1: Group students and their GPAs by dorm

```
mysql> select dorm.name as dorm_name, student.name as  
        student_name, gpa  
        -> from dorm join student using (dorm_id)  
        -> group by dorm_name, student_name, gpa;
```

dorm_name	student_name	gpa
Armstrong	Alice	3.6
Armstrong	Bob	2.7
Armstrong	Cheng	3.9
Brown	Dhruv	3.4
Brown	Ellie	4
Brown	Fong	2.3
Caldwell	Gerd	4
Caldwell	Hal	2.2
Caldwell	Isaac	2
Caldwell	Jacque	5

10 rows in set (0.00 sec)

## Step 2: Get the average GPAs of each dorm.

```
mysql> select dorm.name as dorm_name, avg(gpa) as  
        average_gpa  
        -> from dorm join student using (dorm_id)  
        -> group by dorm_name;
```

```
+-----+-----+  
| dorm_name | average_gpa      |  
+-----+-----+  
| Armstrong | 3.400000015894572 |  
| Brown     | 3.2333333492279053 |  
| Caldwell  | 3.300000011920929 |  
+-----+-----+  
3 rows in set (0.00 sec)
```

## Step 2.1 Formatting Numeric Values

```
mysql> select dorm.name as dorm_name, format(avg(gpa), 2)
      as average_gpa
      -> from dorm join student using (dorm_id)
      -> group by dorm_name;
```

```
+-----+-----+
| dorm_name | average_gpa |
+-----+-----+
| Armstrong | 3.40        |
| Brown     | 3.23        |
| Caldwell  | 3.30        |
+-----+-----+
3 rows in set (0.01 sec)
```

## FORMAT(x,d[,locale])

- ▶ Formats the number *x* to *d* decimals using a format like 'nn,nnn.nnn' and returns the result as a string. If *d* is 0, the result has no decimal point or fractional part.
- ▶ *locale* defaults to the value of the *lc\_time\_names* system variable.

```
mysql> select @@lc_time_names;  
+-----+  
| @@lc_time_names |  
+-----+  
| en_US           |  
+-----+  
1 row in set (0.00 sec)
```

## Step 3: Get max average gpa from average gpa results.

Using a nested query:

```
mysql> select dorm_name, max(average_gpa) as  
      max_average_gpa  
      -> from (select dorm.name as dorm_name,  
                    format(avg(gpa), 2) as average_gpa  
      ->      from dorm join student using (dorm_id)  
      ->      group by dorm_name) as avg_gpas;
```

```
+-----+-----+  
| dorm_name | max_average_gpa |  
+-----+-----+  
| Armstrong | 3.40           |  
+-----+-----+  
1 row in set (0.00 sec)
```



# Views

```
mysql> create view cacm_issues as
-> select * from book
-> where book_title = 'CACM';
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> show tables;
+-----+
| Tables_in_pubs |
+-----+
| author          |
| author_pub      |
| book            |
| cacm_issues     |
| pub            |
+-----+
5 rows in set (0.00 sec)
```

# A View is Like a Table

```
mysql> select * from cacm_issues;
```

book_id	book_title	month	year	editor
1	CACM	April	1960	8
2	CACM	July	1974	8

2 rows in set (0.00 sec)